

Soma Heritage

Author: Dalin Jia (362118251@qq.com)

Affiliation: Independent Researcher

Email: 362118251@qq.com

Date: June 2026

Version: v3.0

Soma Heritage

Soma Heritage: Neural Network Distillation via Signal Field Resonance

Abstract

1. Introduction

1.1

1.2 Soma Heritage

1.3

2. Method

2.1

- W_Q, W_K, W_V, W_O
- $W_c \in \mathbb{R}^{n_{kv} \times k \times d_{head}}$
- $\log \gamma \in \mathbb{R}^k$

2.2

$$\mathcal{L}_{\text{total}} = w_1(t) \cdot \mathcal{L}_{\text{feature}} + w_2(t) \cdot \mathcal{L}_{\text{logit}} + w_3(t) \cdot \mathcal{L}_{\text{consistency}}$$

2.3

```
: M_T, M_S, {L_1, L_2, ..., L_m}
: M_S'

1: for stage = 1 to m do
2:   l ← L_stage
3:   M_S.layers[l].attention ← SignalFieldAttention()
4:   for i = 0 to l do
5:     freeze(M_S.layers[i])
6:   end for
7:   train M_S.layers[l] with three-layer loss
8:   freeze(M_S.layers[l])
9: end for
```

2.4

$$I(l) = \kappa(A_l) \cdot |\nabla \mathcal{L}_l|_2$$

3. Experiments

3.1

3.2

3.3 PPL

Baseline PPL	SignalField PPL			
Layer 0	22.375	23.062	**+3.07%**	
Layer 11	22.375	22.255	**−0.57%**	
Layer 23	22.375	20.011	**−10.57%**	

3.4

3.5

Baseline PPL	SignalField PPL		
WikiText-2	22.375	23.062	+3.07%
Penn Treebank	23.500	22.800	-2.98%

3.6

3.7 FLOPs

** 9SFA vs Standard Attention FLOPs**

SFA	Standard Attention		
FLOPs (seq=1024, d=512)	1.08×10^9	1.61×10^9	** -32.8% **

3.8

4. Discussion

4.1

4.2

4.3

5. Conclusion

References

- [1] Hinton G, Vinyals O, Dean J. Distilling the Knowledge in a Neural Network. *arXiv:1503.02531*, 2015.
- [2] Jiao X, et al. TinyBERT: Distilling BERT for Natural Language Understanding. *EMNLP*, 2020.
- [3] Sanh V, et al. DistilBERT, a distilled version of BERT. *arXiv:1910.01108*, 2019.
- [4] Sun S, et al. ALBERT: A Lite BERT. *ICLR*, 2020.
- [5] Chen S, et al. GradNorm: Gradient Modulation for Equalizing Loss in Multitask Learning. *ICML*, 2018.
- [6] Tishby N, Pereira FC, Bialek W. The Information Bottleneck Method. *arXiv:physics/0004057*, 1999.

[7] Liu M, et al. A Survey on Knowledge Distillation of Large Language Models. *arXiv:2402.13116*, 2024.